

SANJAI KUMAR CHANDRA MOHAN

sanjai.c.kumar@gmail.com | 562-314-2882 | [linkedin.com/in/sanjai-kumar-us/](https://www.linkedin.com/in/sanjai-kumar-us/) | github.com/SanjaiKumar20S115452

EXPERIENCE

Generative AI Scientist Intern

San Francisco, CA

Sephora USA, Inc. · Full-time (on-site)

May 2026 - Sep 2026

- Architected and built **MiRA end-to-end, a production multimodal GenAI platform for Sephora Beauty Chat**, combining real-time voice, image, and text intelligence through Python, LangGraph, LiteLLM, RAG, WebSockets, and PostgreSQL/pgvector for Sephora's **10M+ user digital ecosystem**
- Engineered the **agentic and multimodal AI backend**, integrating STT/TTS, LLM routing, vector retrieval, catalog grounding, and computer vision while benchmarking **5+ speech models** and indexing **12K+ products** for retrieval
- Owned MiRA from **AI architecture through Azure/Kubernetes deployment and production troubleshooting**, diagnosing KEDA autoscaling, gateway, and model-routing failures through log analysis while presenting the system and its results to engineering and leadership stakeholders

AI Automation Engineer - Agentic Systems & Orchestration

San Jose, CA

Wordcorn AI · Full-time (hybrid)

Jan 2022 - Dec 2025

- Built an **AI-powered automation system** that generated, scheduled, and published content across multiple channels, executing **30+ automated workflows weekly with zero human intervention for 9 months**
- Engineered **Python automation pipelines and API integrations** for long-running workflows and duplicate prevention
- Built end-to-end testing and debugging workflows using **Selenium and Postman**, reducing recurring failures by 35%

PROJECTS

MiRA (Multimodal Intelligence and Reasoning Architecture) | Python, LangGraph, LiteLLM, CLIP, LoRA, PostgreSQL (pgvector)

- Built a **production Agentic AI and multimodal system** that understands voice, image, and text, orchestrates LLMs and speech models, and retrieves products across **12K+ catalog items with approximately 200 ms retrieval latency**
- Designed the distributed architecture around **LangGraph agents, LiteLLM model routing, vector search, WebSocket streaming, PostgreSQL/pgvector, Docker, Kubernetes, and Azure** for scalable real-time AI execution

AI Workflow Automation & Scheduling Platform | FastAPI, Celery, Redis, LangGraph, Docker, Kafka, Kubernetes, GCP

- Built an **AI-enabled workflow orchestration platform** for asynchronous job execution, scheduling, and multi-channel actions, reducing manual coordination by approximately **80%**
- Architected **distributed processing** using FastAPI, Celery, Redis, PostgreSQL, Kafka-style event pipelines, Docker

TECHNICAL SKILLS

- **Languages & Scripting:** Python, Java, JavaScript, TypeScript, C, C++, SQL, Bash
- **Frameworks & Libraries:** FastAPI, Celery, Redis, React, Node.js, LangGraph, LangChain, Selenium
- **AI/ML:** Generative AI, LLMs, Agentic AI, Prompt Engineering, RAG, Multimodal AI, NLP, Computer Vision, PyTorch
- **Data & Cloud:** AWS, Azure, GCP, Kubernetes, Docker, PostgreSQL, MySQL, NoSQL, Redis, Kafka, Git
- **Software Engineering:** Backend Engineering, Infrastructure Engineering, Distributed Systems, REST APIs

EDUCATION

California State University, Long Beach

Long Beach, California

M.S. Computer Science, Concentration in Machine Learning (GPA: 3.87)

Expected, December 2026